

Agente Crédito LATAM

WORKFLOW 4 · CREDIT-PRODUCT INFORMATION AND ELIGIBILITY

A credit assistant a bank can switch on.

The model talks, rules decide, every figure is cited, and it knows when not to act. In Spanish and Portuguese, for Colombia, Mexico and Argentina.



Eladio Yovera

Data platform

Medallion pipelines and demand insights



Manuela Larrea

AI / ML

Agent, rules, retrieval, pre-score, evals



Nicolle Mayol

Service

Azure, CI/CD, load tests, front end

TRY IT

Live app

gentle-moss-06f3f1f0f.2.azurestaticapps.net


Repository

github.com/NicolleMayol/factored-hackathon-2026-the-trident

How to navigate

- 1 Home: the numbers and how it decides.
- 2 “See the customer experience”, then pick a test customer.
- 3 Ask for a rate, run a pre-check, then try “ignore your rules”.

The most repeatable contact, and where an AI must never decide



PRODUCT CONTACTS

22%

150,863 contacts · FCR 0.896

MX 50% CO 30% AR 20%

ABSORB

Absorb the repeatable

Solved on first contact 9 times in 10: the volume an assistant should take.

NEVER DECIDE

Never decide the edge

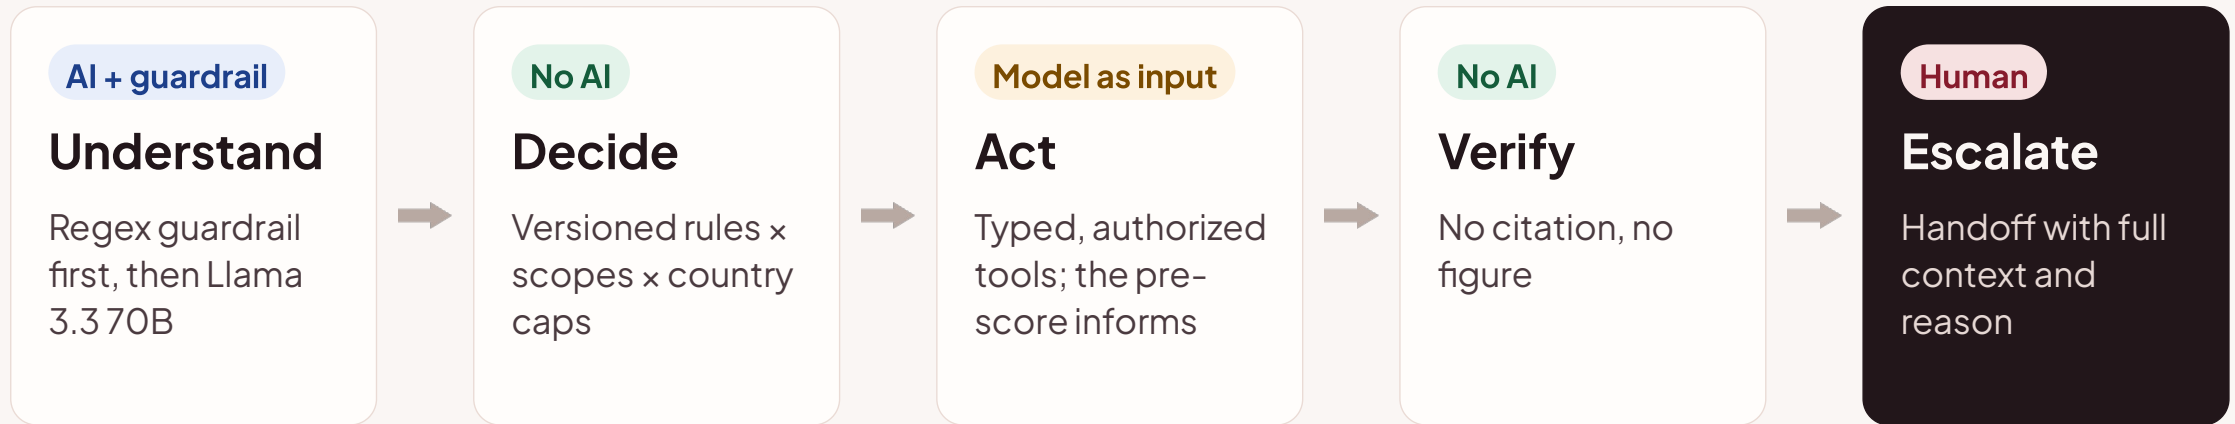
48% of Basic customers sit near the score cutoff; 15% have no score, 20% no income.

OUR CHOICE · WORKFLOW 4

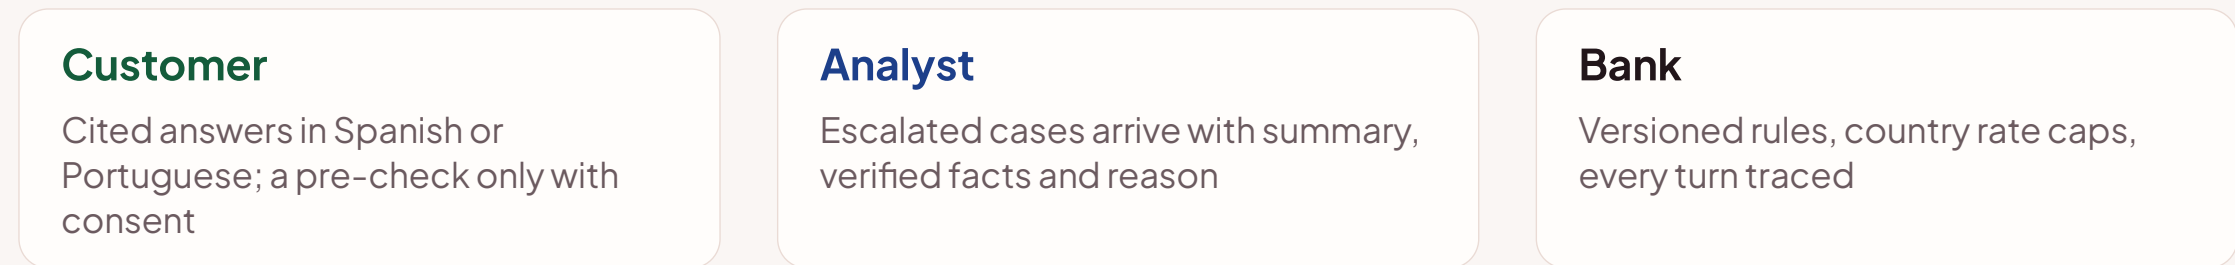
It knows when not to act

Cited answers, a pre-check only with consent, and a person for every edge case.

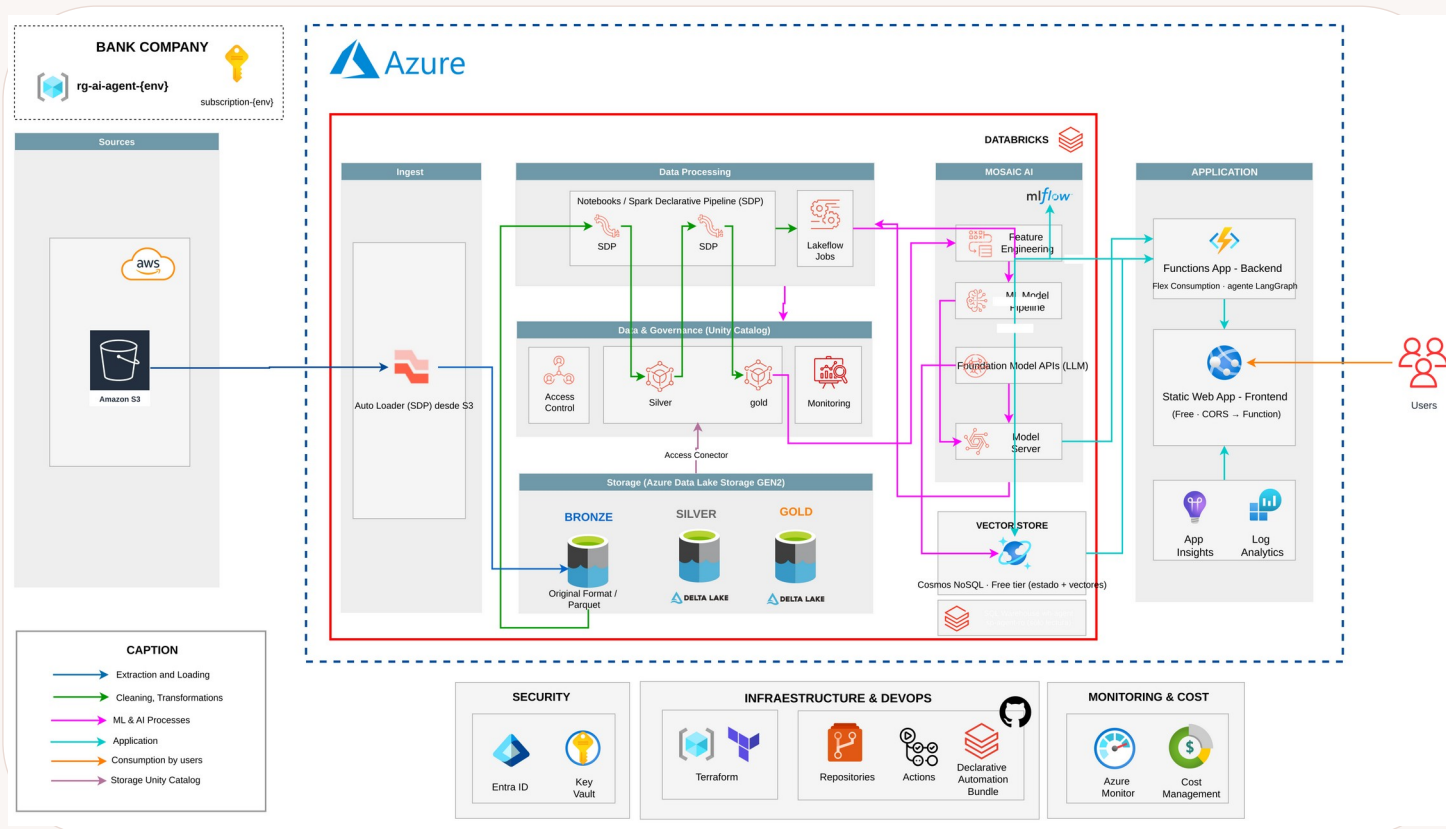
The model talks. Rules decide. Every figure is cited.



WHAT THE BANK SWITCHES ON · WHO GETS WHAT



Built the way a bank runs it, all in code



DESIGN CHOICES A BANK CARES ABOUT

RAG · ES + PT
Answers only from bank policy
Hybrid search in Cosmos DB, bge-m3 embeddings.
Recall@5: 1.0 es, 0.92 pt.

EXTERNAL SOURCES
Regulators as data, cited
SFC, Banxico and BCRA rate caps with URL and date.
Attacks from MITRE ATLAS.

ORCHESTRATION
A fixed graph, not a free agent
LangGraph: the LLM never approves, versioned rules do. Same input, same decision, auditable.

TO-BE · IN A BANK
Private endpoints, WAF, Entra ID, hub-and-spoke

Open the to-be diagram ↗

1aC Terraform + Actions, approval gates	8 contracts first: 3 roles in parallel	29 ADRs, one per decision	1 own agent skill for docs as code	210 tests + eval on every PR
--	---	-------------------------------------	---	--

It knows when not to act, on real data

0.98

Abstain accuracy

0.019

False action rate

0

Successful prompt injections

Injections · deterministic guardrail



Intent failures · prompt v1 → v2



Pre-score AUC · portfolio features



Live, all-real, and built to be operated.

Dependencies

5/5 ok

live /api/healthz

p50, 10 users

2.7 s

89 turns in 1 min

p95, 10 users

17.5 s

target 8 s, missed

Cost

\$0.00035

per turn

Tracing

Every turn

ops.agent_turns

Underload

84 of 89

answered; 5 safe
handoffs, never a broken
answer

What's next, and why

p95 under load → provisioned LLM throughput and retry on rate limits

Real intent labels → label real call transcripts

CO and MX rate caps → official monthly figures, cited

Seconds, one axis

